

Harrison Kinsley

Harrison@pythonprogramming.net

Word Count: 41532 (tools)

Links to the chapters:

[Chapter 1 - Introducing Neural Networks](#)

[Chapter 2 - Coding Our First Neurons](#)

[Chapter 3 - Adding Layers](#)

[Chapter 4 - Activation Functions](#)

[Chapter 5 - Loss](#)

[Chapter 6 - Optimization](#)

[Chapter 7 - Derivatives](#)

[Chapter 8 - Gradients, Partial Derivatives, and the Chain Rule](#)

[Chapter 9 - Backpropagation](#)

[Chapter 10 - Optimizers](#)

[Chapter 11 - Testing Data](#)

[Chapter 12 - Validation Data](#)

[Chapter 13 - Training Dataset](#)

[Chapter 14 - L1 and L2 Regularization](#)

[Chapter 15 - Dropout](#)

[Chapter 16 - Binary Logistic Regression](#)

[Chapter 17 - Regression](#)

[Chapter 18 - Model Object](#)

[Chapter 19 - A Real Dataset](#)

[Chapter 20 - Model Evaluation](#)

[Chapter 21 - Saving and Loading Model Information](#)

[Chapter 22 - Model Predicting/Inference](#)

If you've not joined the Google group yet, you will need to accept the invite to view the individual chapters above. Contact me @ Harrison@pythonprogramming.net if you're having issues.

you can also go file>download or file>make a copy for a private version.

